

Murphi Final Report

timdunn@umich.edu

Approach

I started this assignment by reading the existing IV protocol in Murphi until I fully understood how it worked. Next, I began working from that code base and fully replaced the IV protocol with the MSI protocol from the textbook, which assumes network ordering [1].

My initial approach was to repeatedly iterate the following steps: compile, run Murphi, detect unsatisfied invariants, and tweak the existing protocol to avoid the new issue introduced by random-order packet delivery. This process was time-consuming and difficult. After many hours of work and little tangible progress, I eventually gave up on this approach and decided to redesign the whole protocol at once from a higher level.

While spending hours fixing individual coherency protocol flaws, I realized that all my issues were the result of state inconsistencies between the **Directory** and **Processor**, stemming from the fact that the **Directory** assumed the **Processor** had received packets which it hadn't yet, and that the **Processor** had left a transient state when it in fact hadn't. To redesign the protocol, I wrote out the simplified protocol from the textbook, and then slowly worked my way through each possible state, adding **Directory** transient states and acknowledgements from the **Processor** to **Directory** wherever necessary. After reworking my high-level design a few times and debugging, the MSI protocol worked! From there, extending unordered MSI to unordered MESI was fairly straightforward.

Protocol

Messages for my MSI protocol implementation were sent over four virtual channels. The table below shows the organization of message types into channels such that for each possible sequence of messages and responses, a later response will not be passed along the same channel as an earlier message, which would introduce the possibility of deadlock.

RequestChannel	ForwardChannel	ResponseChannel	OrderChannel
GetS	FwdGetS	Data	DataAck
GetM	FwdGetM, Inv	Data, InvAck	DataAck
PutS		PutAck	PutAckAck
PutM		PutAck	PutAckAck

As the table above shows, my protocol introduces **DataAck** and **PutAckAck** messages, which were not necessary for the protocol described in [1], which assumes an in-order network. Since for this assignment we could not assume that messages would be received in the order in which they were sent, the simplest approach to ensure coherence was to introduce several transient states for both the **Processors** and the **Directory**, and use the new acknowledgment message types to regulate movement from one state to another.

Details are shown in the attached full state diagrams. My protocol may involve unnecessary stalling, but correctness (not efficiency) is the stated goal of this project. Note that although

a chain of four messages are sent when a **Processor** requests permission to read/write a new memory location (e.g. **GetS**→**FwdGetS**→**Data**→**DataAck**), it is still a 3-hop protocol because the **Processor** receives new data after the third hop, and the **Data** message is passed directly from the **Owner** to the **Processor** without first travelling through the **Directory**.

In addition to each **Processor** storing a value, **Processors** also needed to keep track of the number of **Ack** requests they were expecting and had received for certain states. This was done by introducing **acks_needed** and **acks_received** variables. Messages were extended beyond containing information regarding the **mtype**, **dst**, **src**, **vc**, and **val**, to also include a **who** and **ack**. These additions told processors where to send data for **FwdGetS**/**FwdGetM** requests and how many acknowledgements the **Processor** should wait for.

Optimization

The optimization I chose to implement for this project was adding an **E** state and implementing an **MESI** protocol. This required the introduction of **PutE** and **EData** messages (which were placed in the **RequestChannel** and **ResponseChannel**, respectively), as well a few new **Processor** (**E**) and **Directory** (**E**, **E.P**, **E.A**) states. I really learned a lot from this project, but felt that there would be no additional benefit in struggling to ensure the correctness of a more “difficult” optimization, since I had already learned a lot from redesigning protocols and investigating coherency protocol edge cases.

References

[1] Sorin, Daniel J., Mark D. Hill, and David A. Wood. "A primer on memory consistency and cache coherence." Synthesis lectures on computer architecture 6.3 (2011): 1-212.

MESI: Processor

	load	store	evict		FwdGetS	FwdGetM	Inv	PutAck	Exclusive Data from Dir	Data from Dir, ack=0	Data from Dir ack>0	Data from Owner	InvAck, ack>0	InvAck, ack=0	DataAck
I	send GetS to Dir / IS_D	send GetM to Dir / IM_AD													
IS_D							stall		send DataAck to Dir / E	send DataAck to Dir / S		send DataAck to Dir + Owner / S			
IM_AD					stall	stall				send DataAck to Dir / M	- / IM_A	send DataAck to Dir + Owner / M	ack--		
IM_A					stall	stall							ack--	send DataAck to Dir / M	
S		send GetM to Dir / SM_AD	send PutS to Dir / SI_A				send InvAck to Req / I								
SM_AD					stall	stall	send InvAck to Req / IM_AD			send DataAck to Dir / M	- / SM_A	send DataAck to Dir + Owner / M	ack--		
SM_A					stall	stall							ack--	send DataAck to Dir / M	
M			send PutM+Data to Dir / MI_A		send Data to Req + Dir / MS_A	send Data to Req / MI_AA									
E		- / M	send PutE to Dir / MI_A		send Data to Req + Dir / MS_A	send Data to Req / MI_AA									
MI_A					send Data to Req + Dir / MI_SI_AA	send Data to Req / MI_II_AA		send PutAckAck to Dir / I							
SI_A							send InvAck to Req / II_A	send PutAckAck to Dir / I							
II_A								send PutAckAck to Dir / I							
MI_AA								stall							- / I
MI_II_AA								stall							- / II_A
MS_A						stall		stall							- / S
MI_SI_AA						stall		stall							- / SI_AA

MESI: Directory

[illegible]